



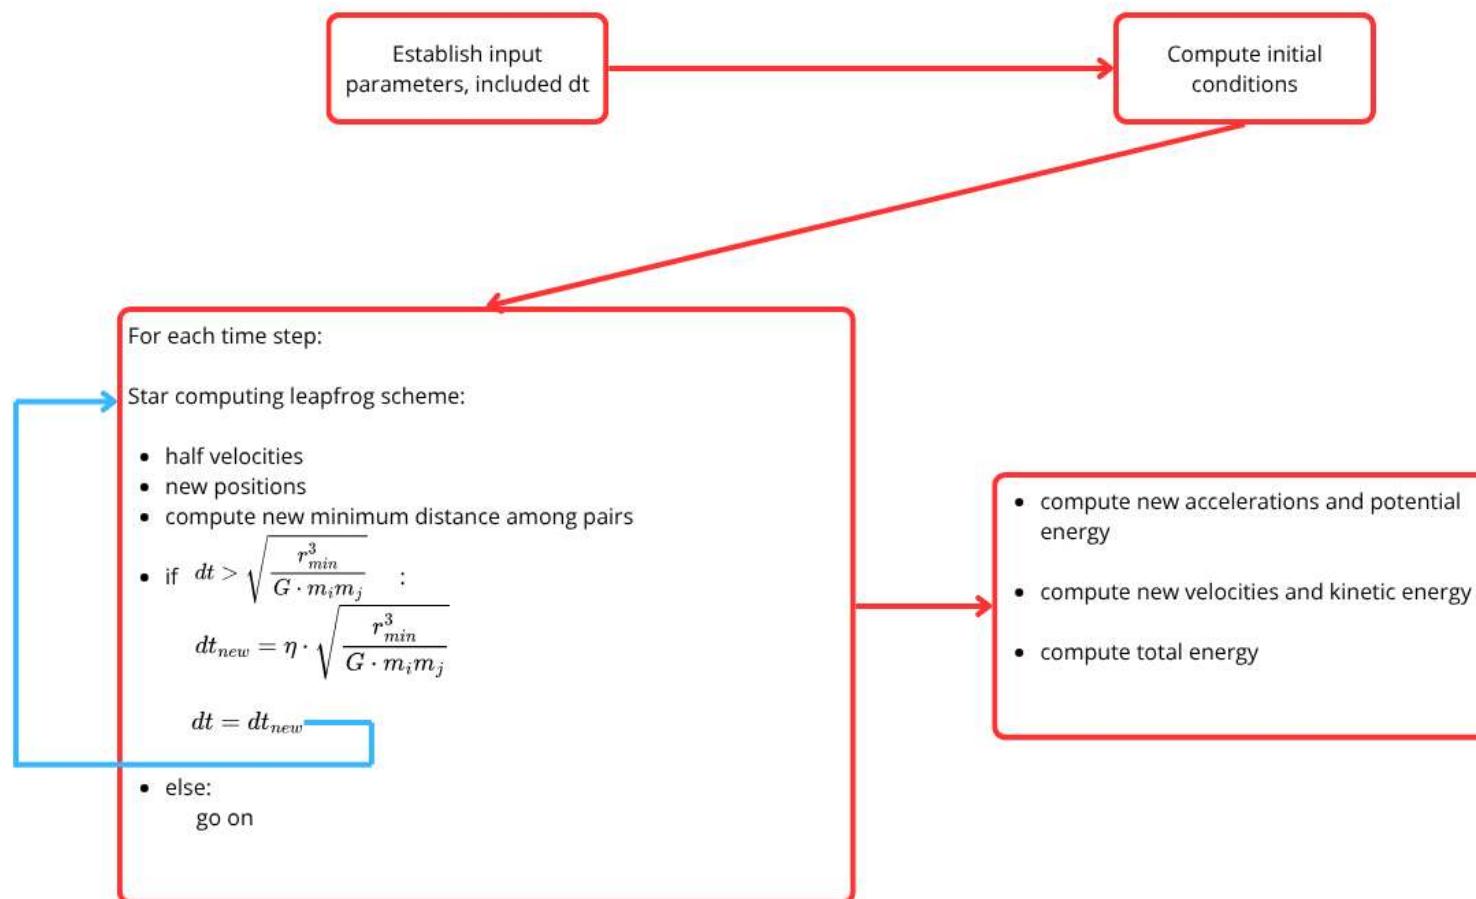
MODERN COMPUTING FOR PHYSICS

N-BODIES SIMULATION WITH
GPU PARALLELIZATION

The n-bodies problem

- We consider a system of n **bodies**, each with a given **mass**, initial **position**, and initial **velocity**
- The bodies interact solely through **gravitational forces**
- We want to simulate their **time evolution**, obtaining their trajectories
- We want to check the **conservation** of total **energy**
- Computational cost **$O(n^2)$**

Serial code in C



Issues: scale of the problem and data types

This scale is chosen such that:

- we can graphically **observe trajectories**: gravitational forces dominate over rectilinear uniform motion) and proportioned in the space
- the numerical scheme is **stable**
- you can compile with both float or **double** precision

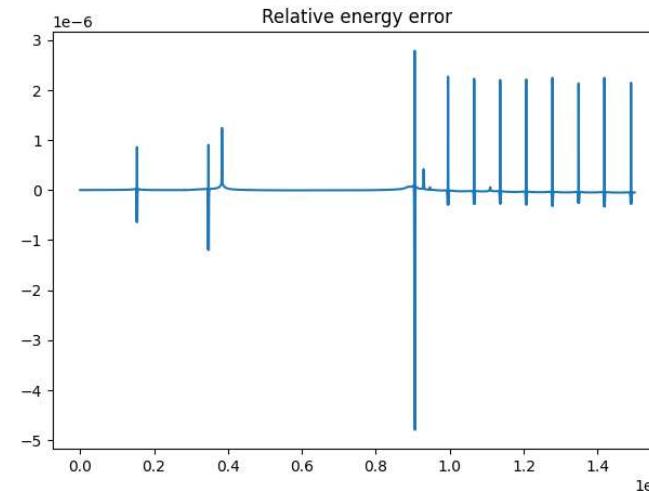
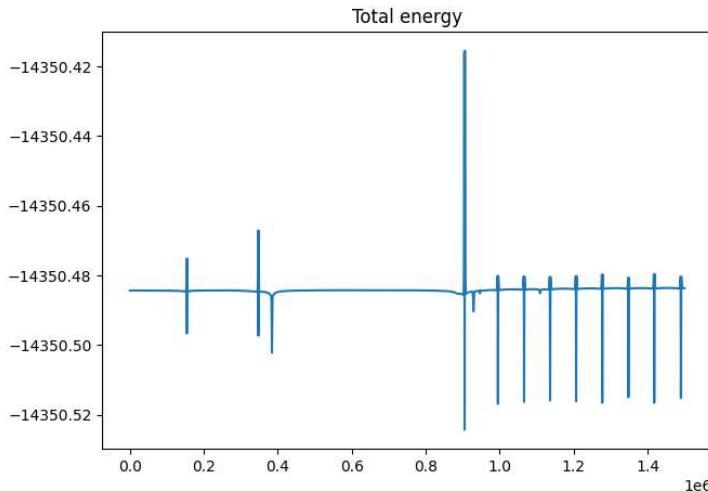
```
for (int i = 0; i < n_bodies; i++) {
    m[i] = rand_interval(1.0f, 10.0f);
    x[i] = rand_interval(0.3f * L, 0.7f * L);
    y[i] = rand_interval(0.3f * L, 0.7f * L);
    z[i] = rand_interval(0.3f * L, 0.7f * L);
    vx[i] = rand_interval(-1.0f, 1.0f);
    vy[i] = rand_interval(-1.0f, 1.0f);
    vz[i] = rand_interval(-1.0f, 1.0f);
}
```

```
#ifdef FLOAT
typedef float real;
#define SORT sqrtf
#define FMOD fmodf
#else
typedef double real;
#define SORT sqrt
#define FMOD fmod
#endif

#define n_bodies 6
#define steps 1500000
#define SRAND 31265

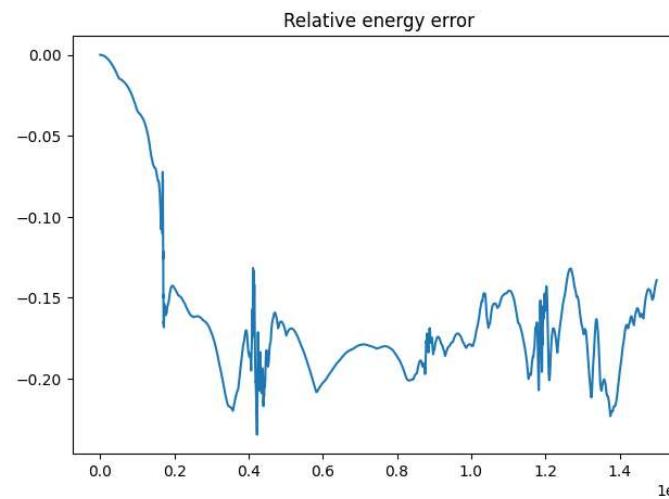
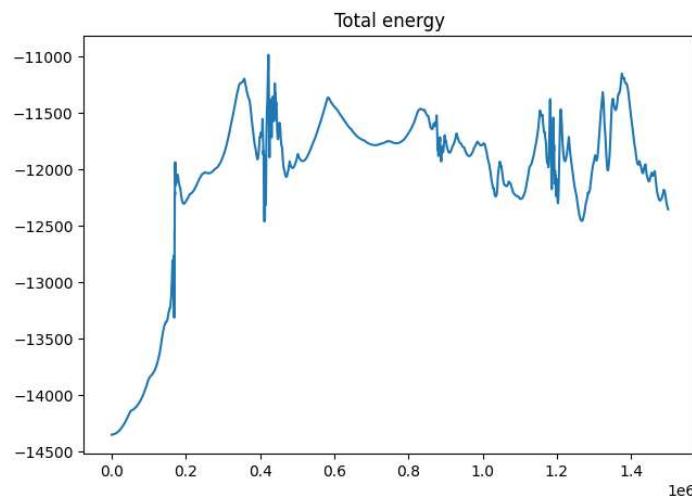
const float dt_max = 1e-6f;
const float eta = 1e-3f;
const float L = 100.0f;
const float G = 1.0e3f;
```

Evolution of total energy of the system



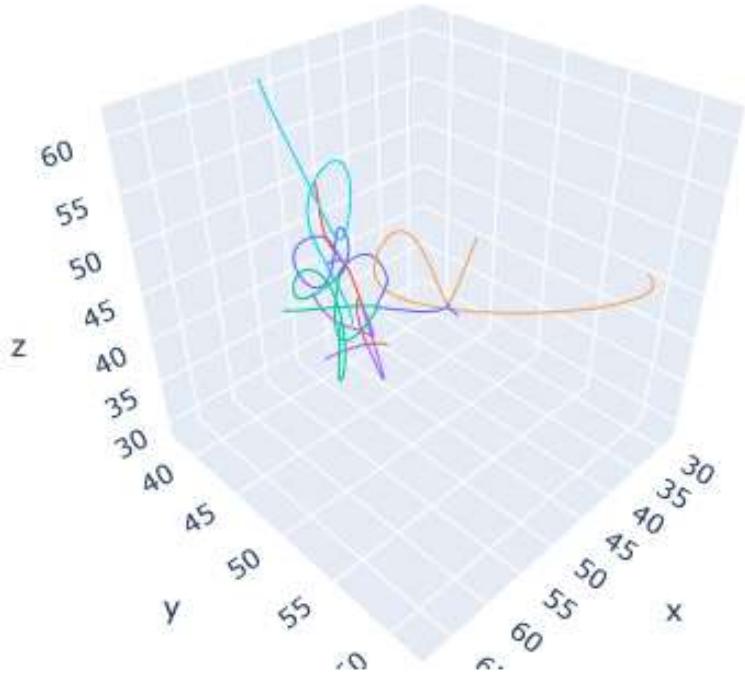
Double precision
relative errors $\sim 10^{-6}$

Relative errors: $\frac{E_t}{E_{t=0}} - 1$



Float precision:
unstable

6 bodies trajectories



64 bodies (with periodic boundary conditions)

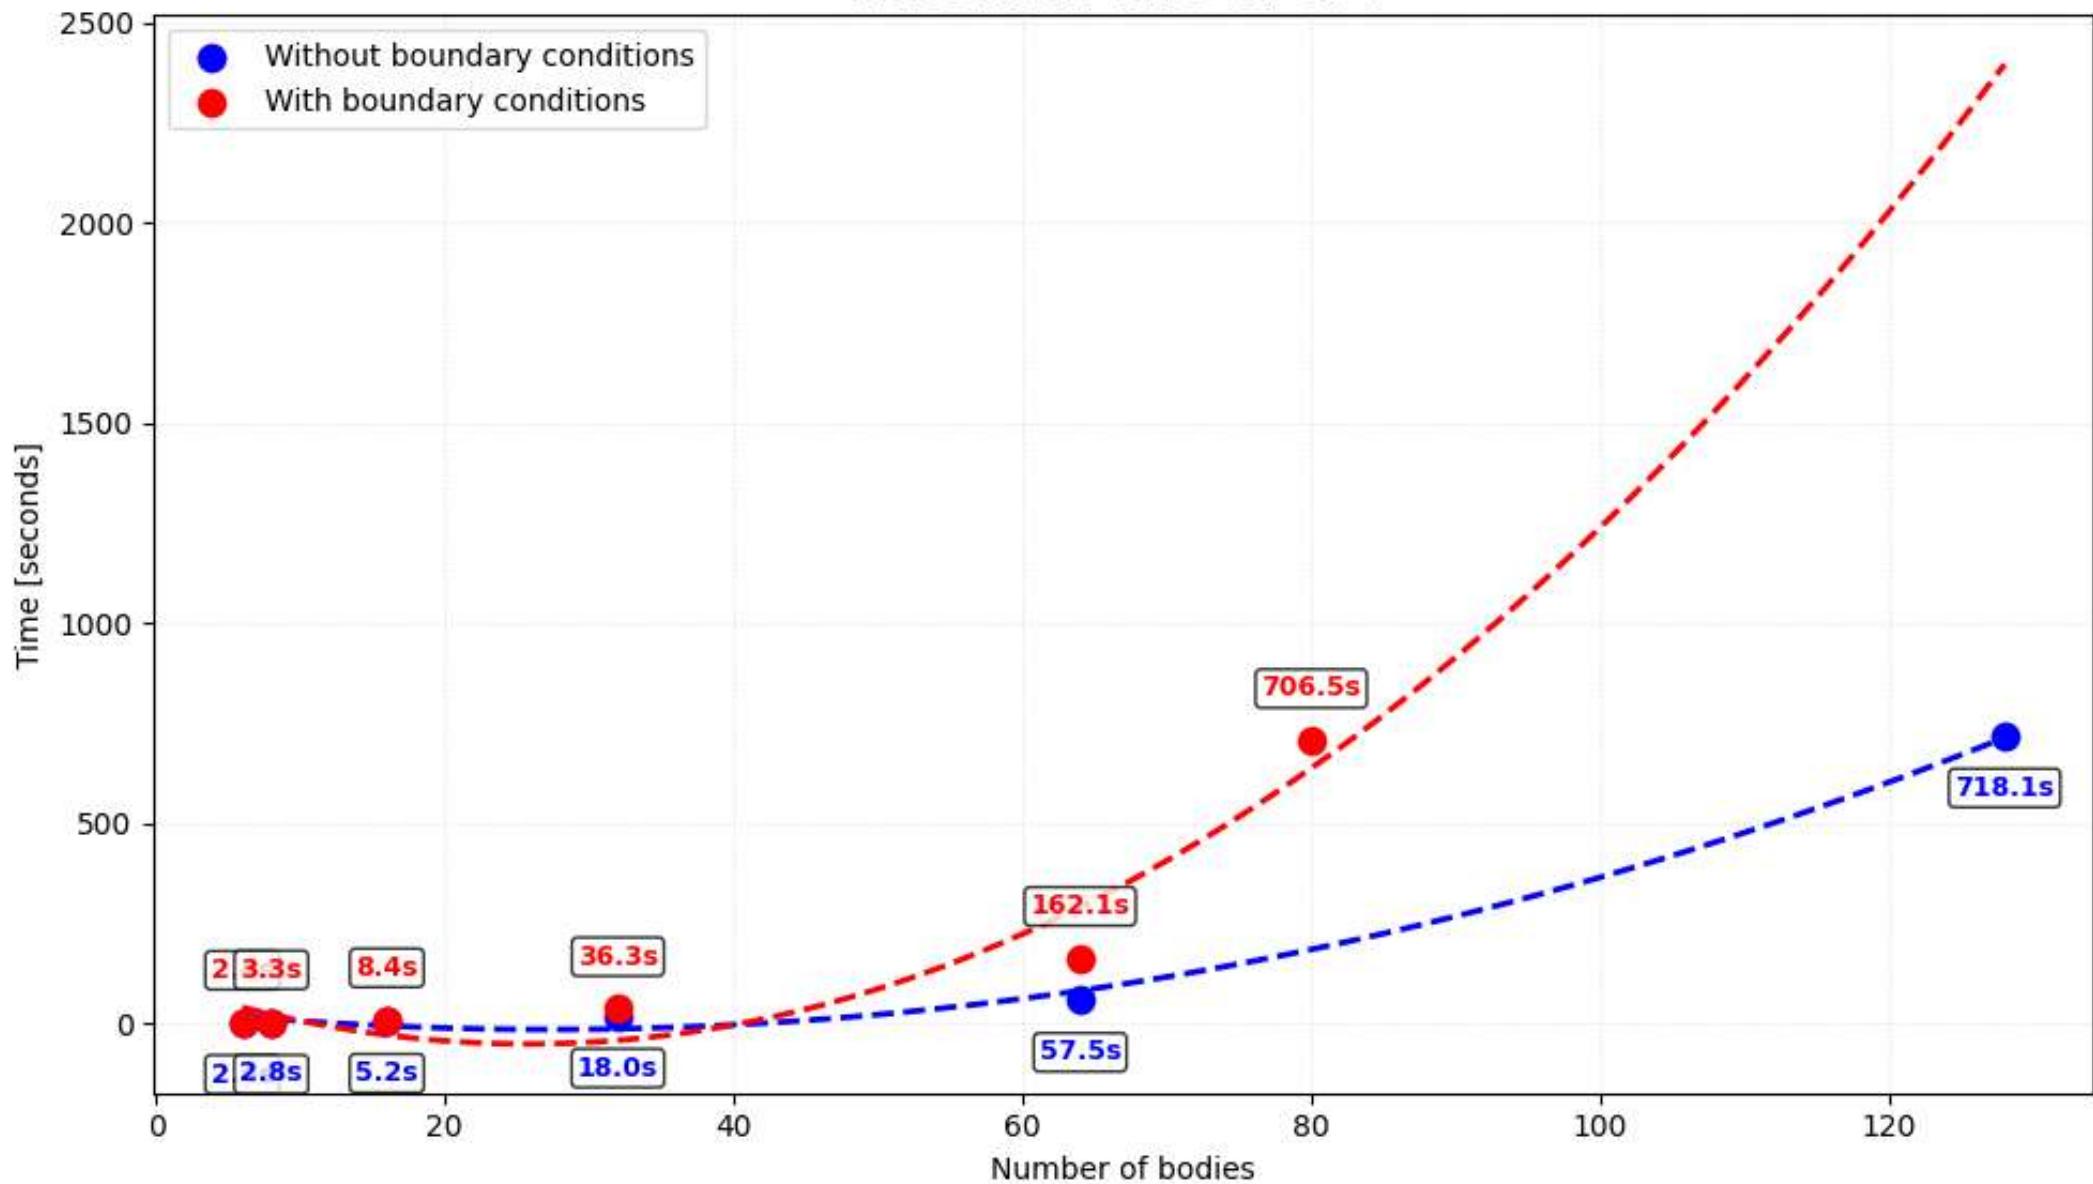


In the code a **csv** file to keep track of the positions has been created as output.

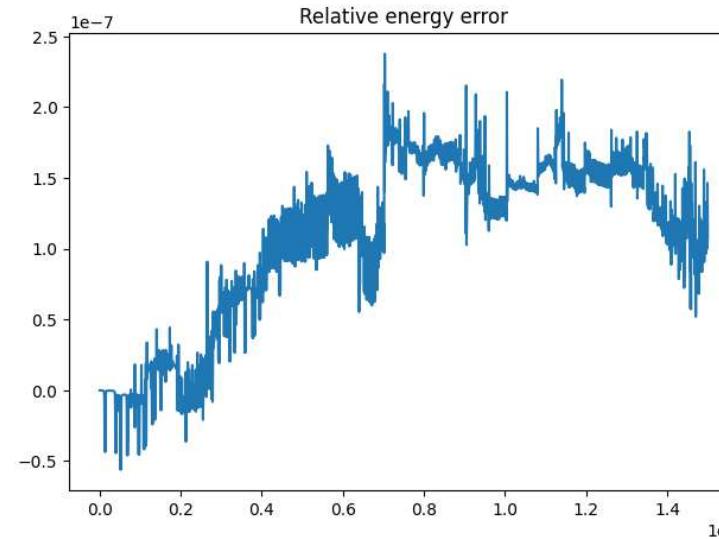
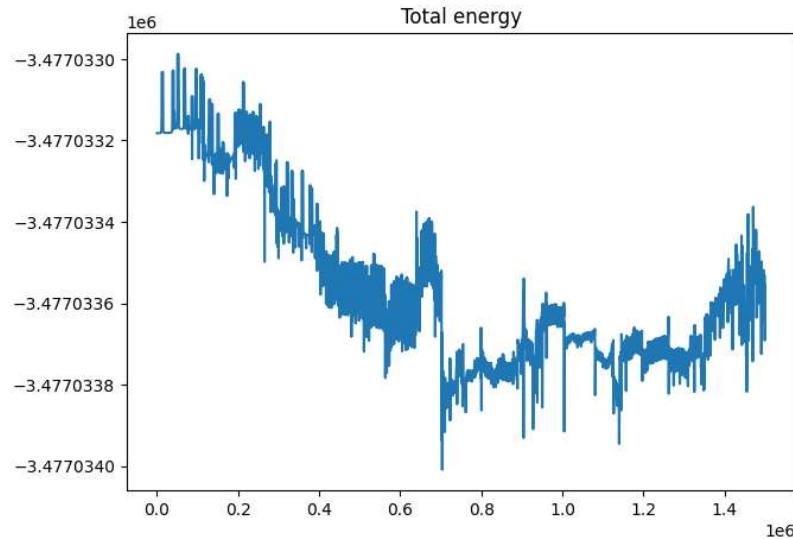
Trajectories have been plotted in python from this file with a script written by generative AI.

The csv has been modified through a script and be animated through OVITO.

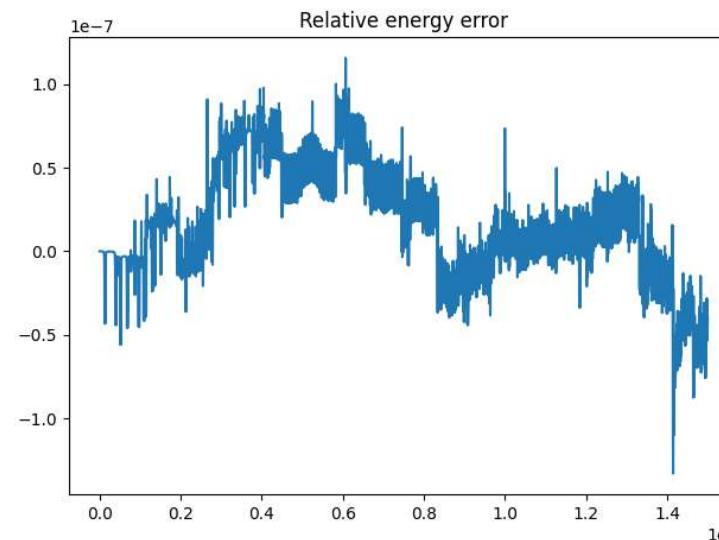
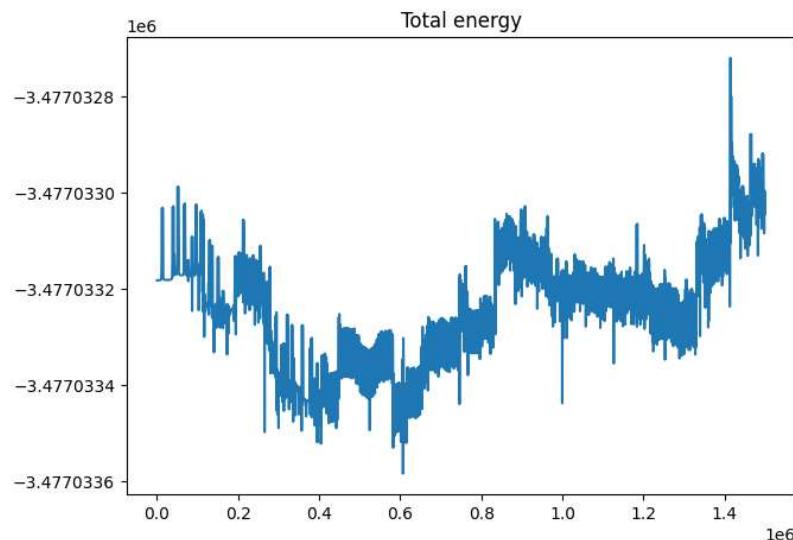
Execution time on CPU



Conservation of energy for a 64 bodies system and 1500k steps



With periodic boundary conditions



Without boundary conditions

Parallel code in CUDA

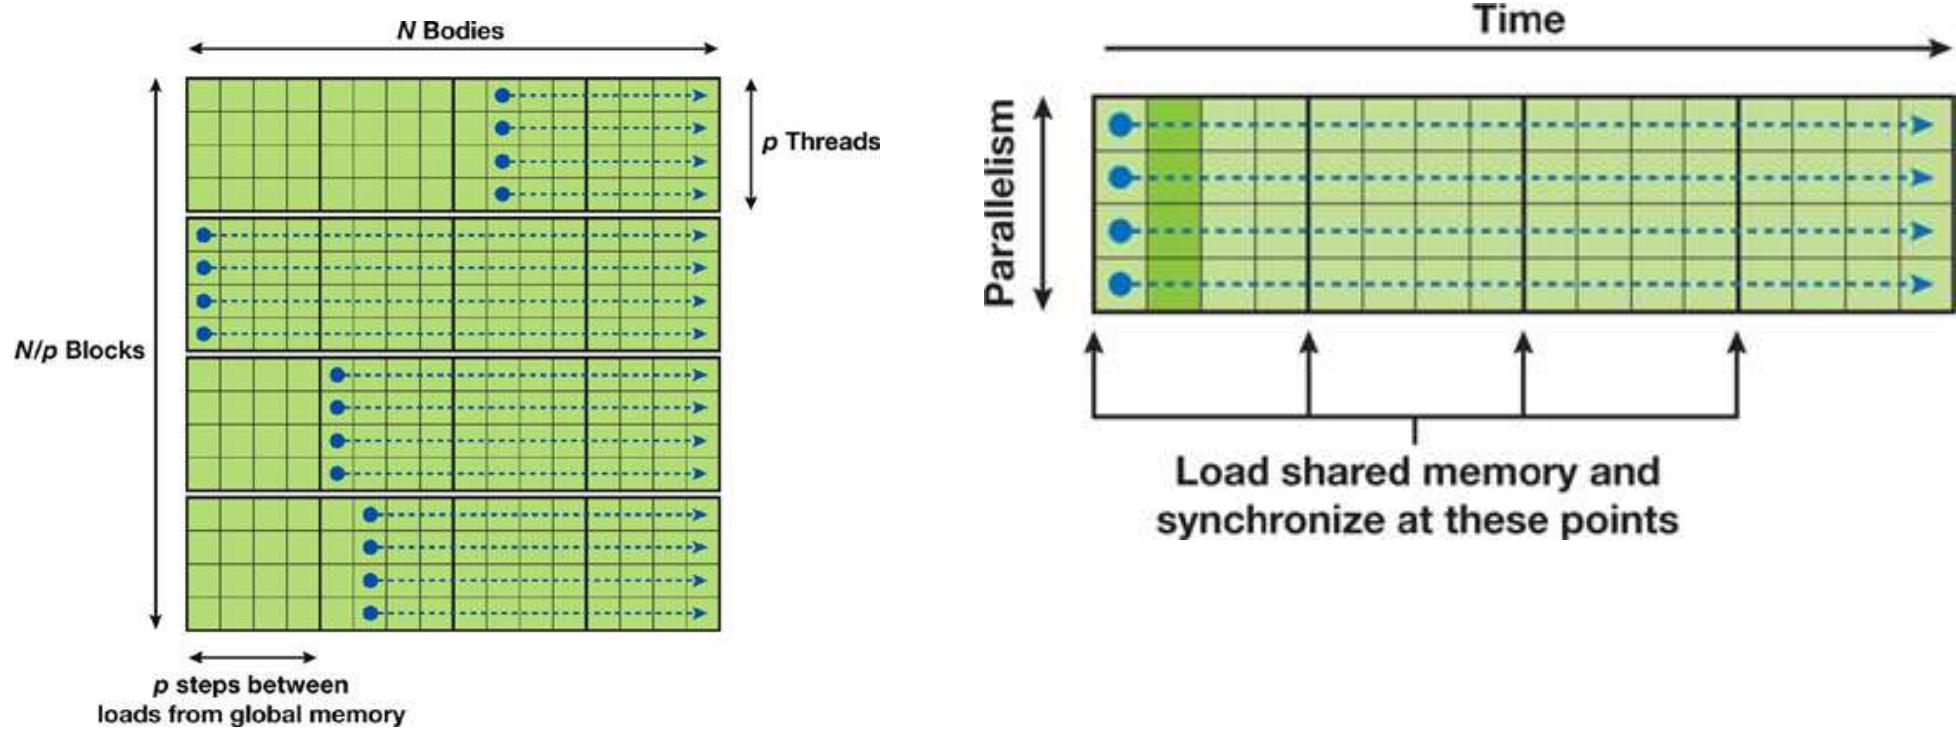
GPU	Tesla T4
Global Memory	15360 MiB
Streaming MultiProcessors	40
Total CUDA Cores	2560
Max Threads per Block	1024
Shared Memory per SM	64 KB
Compute Capability	7.5

Crucial point is updating the accelerations of the bodies

→ **tiling** strategy is used, using the approach suggested in

<https://developer.nvidia.com/gpugems/gpugems3/part-v-physics-simulation/chapter-31-fast-n-body-simulation-cuda>

- We may think the algorithm as calculating each entry \mathbf{a}_{ij} in an $N \times N$ grid of all pair-wise forces.
- Then the total acceleration \mathbf{a}_i on body i is obtained from the sum of all entries in row i .
- $O(N^2)$ would be too large
- We introduce tiles, square regions of the grid $p \times p$
- Boundary conditions not applied



The properties of the bodies are written in shared memory of the blocks for each dimension of the grid.

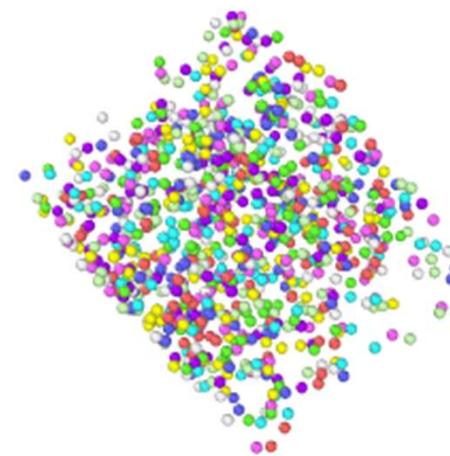
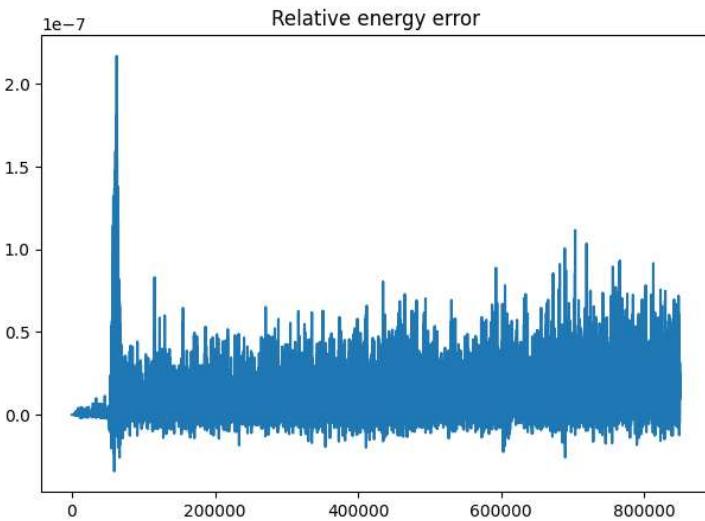
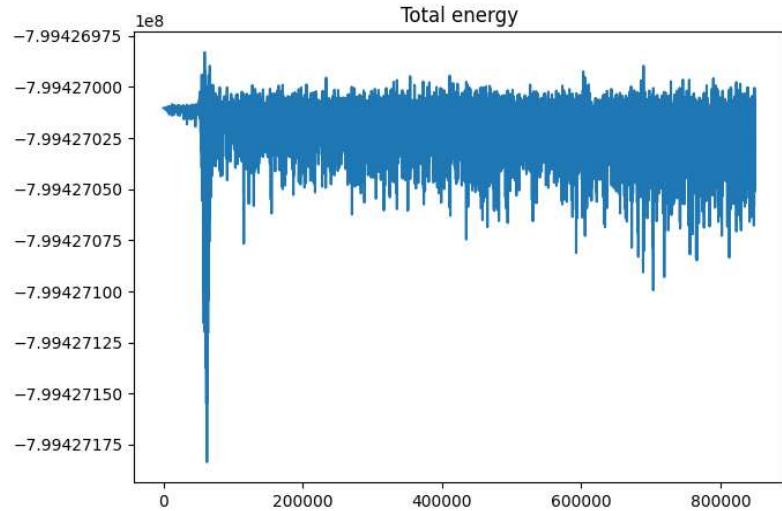
Each **row** corresponds to a **thread** which computes sequentially interactions with the bodies in the columns, summing up accelerations sequentially.

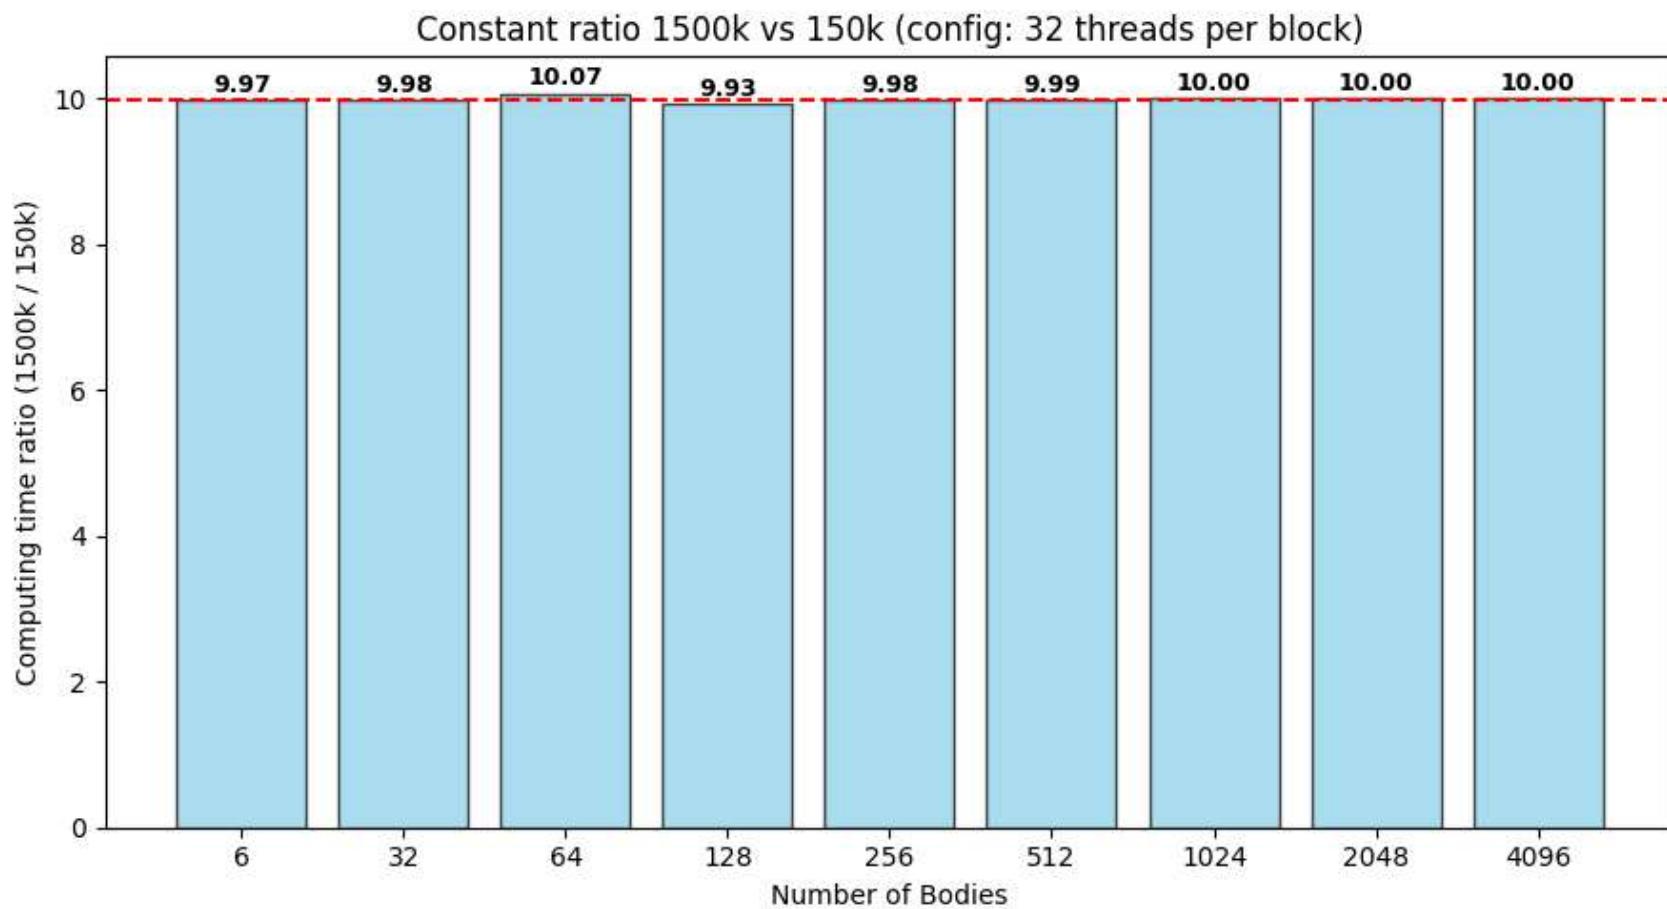
A final sum is computed among all tiles.

The problem's symmetry can't be exploited, the real advantage is the **reduction of accesses to global memory**.

The algorithm preserves conservation of energy also with **1024** bodies

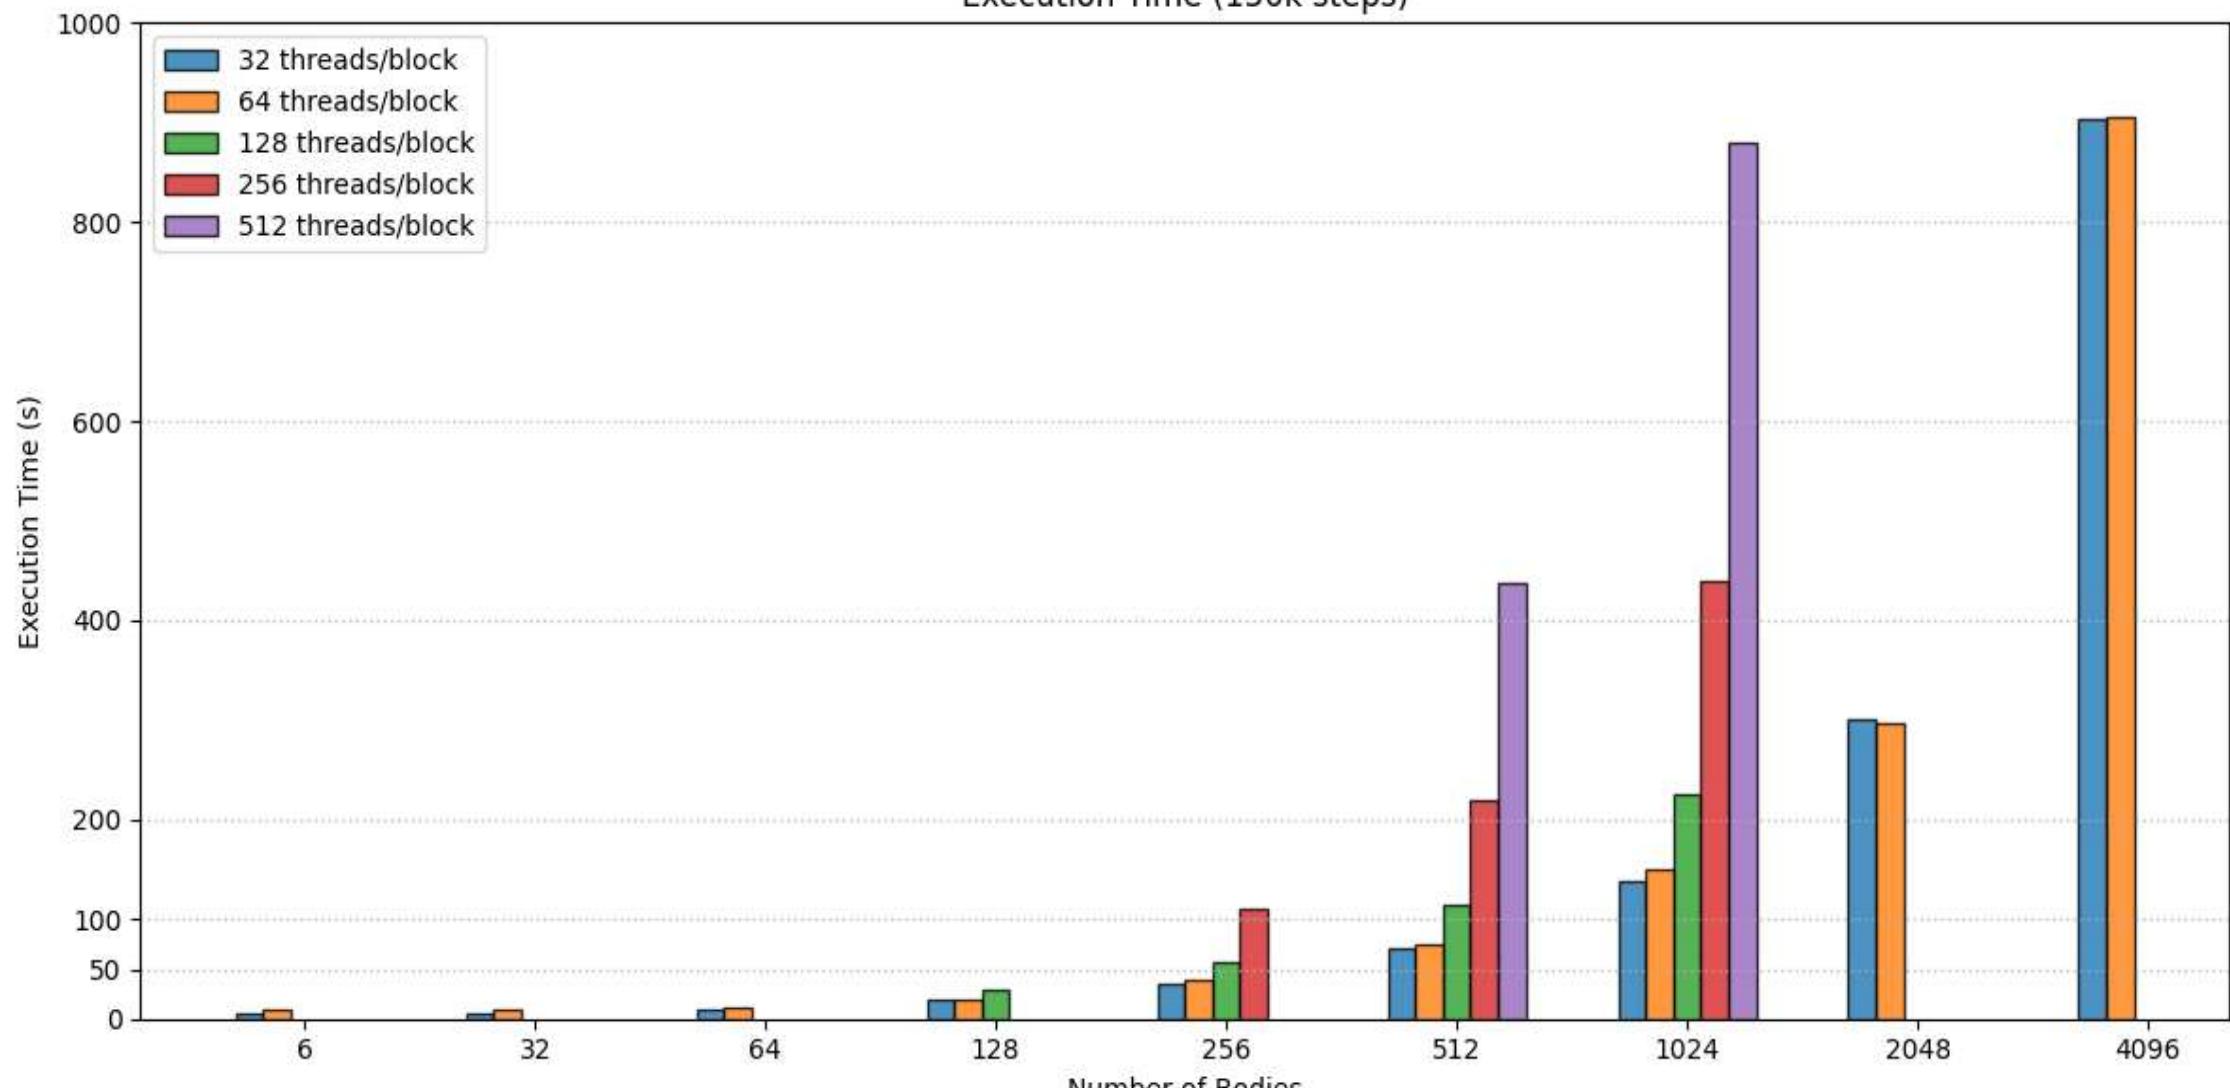
Relative errors $\sim 10^{-7}$





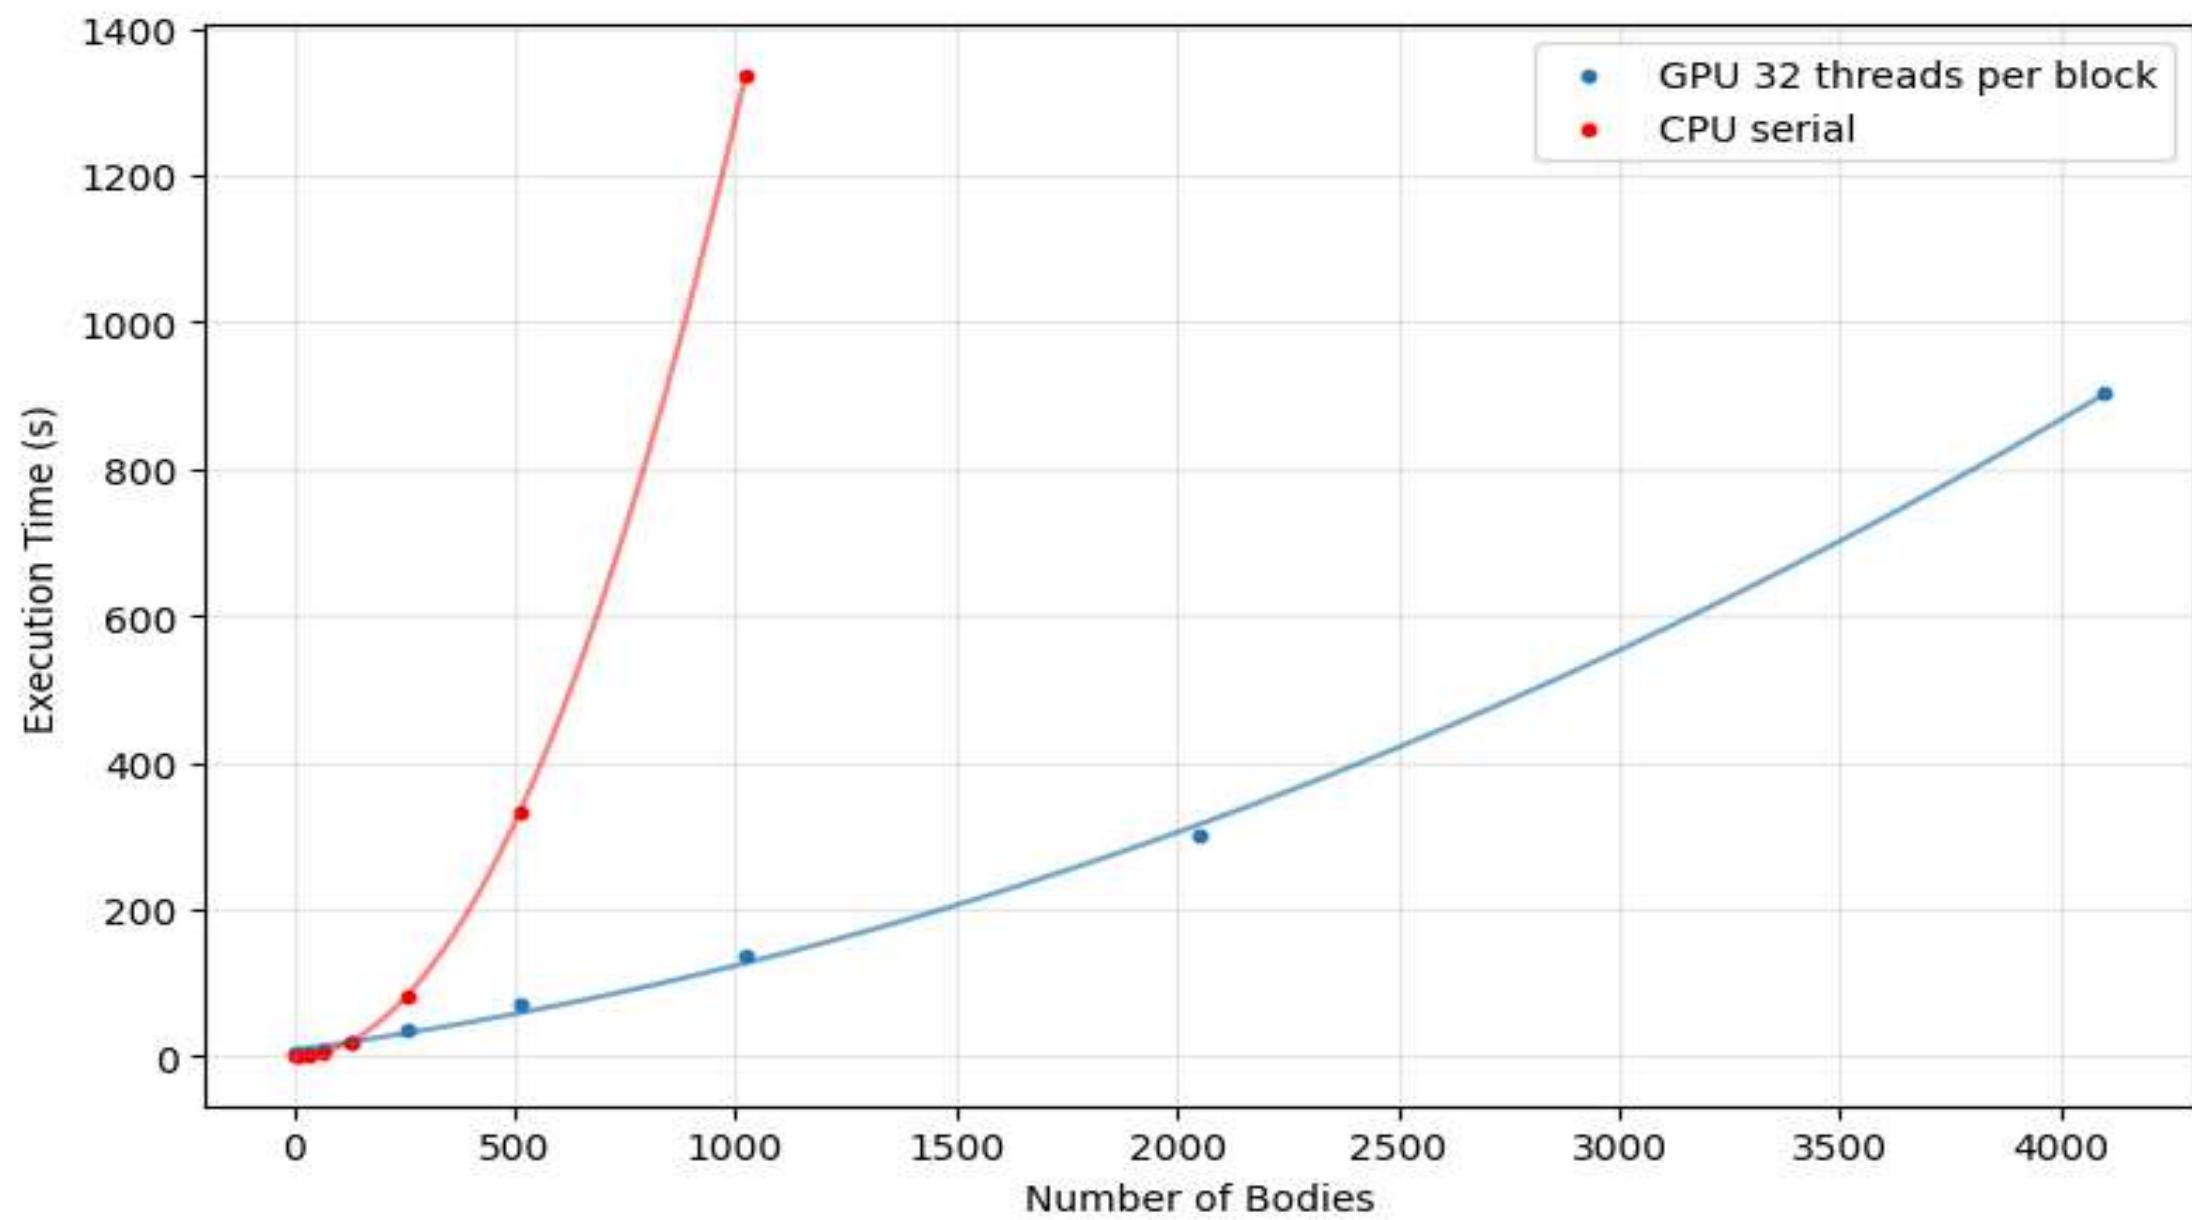
Execution time is linear with respect to the number of steps.

Execution Time (150k steps)



`n_blocks=n_bodies/thr_per_b`

Having more blocks hides latency due to memory transfers.



GPU Execution Time vs Bodies (1500k steps)

